

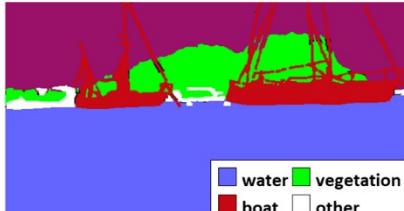
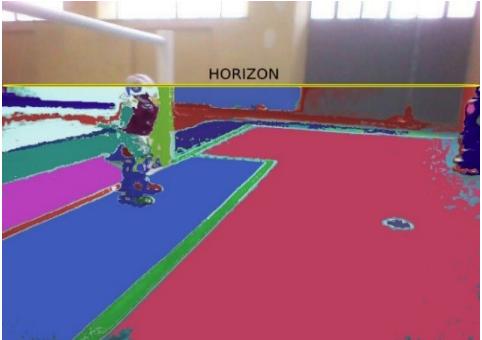


**UNIVERSITÀ DEGLI STUDI  
DELLA BASILICATA**

*Corso di Visione e Percezione*  
A.A. 2019/2020

# Omografie

Aprile 2020



Docente  
**Domenico Daniele Bloisi**

# Il corso

---

- Home page del corso  
<http://web.unibas.it/bloisi/corsi/visione-e-percezione.html>
- Docente: Domenico Daniele Bloisi
- Periodo: Il semestre marzo 2020 – giugno 2020

Martedì 17:00-19:00 (Aula GUGLIELMINI)

Mercoledì 8:30-10:30 (Aula GUGLIELMINI)

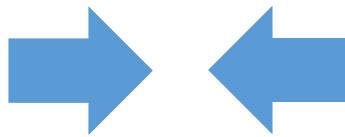
# Riferimenti

---

- Queste slide sono adattate da
  - Noah Snavely - CS5670: Computer Vision  
["Lecture 7: Transformations and warping"](#)
  - M. Brown and D. G. Lowe  
[Recognising Panoramas](#)
- I contenuti fanno riferimento al capitolo 3 del libro "Computer Vision: Algorithms and Applications" di Richard Szeliski, disponibile al seguente indirizzo  
<http://szeliski.org/Book/>

# Image alignment

---



# Image alignment

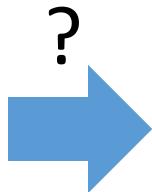
---



# Similarity

---

What is the geometric relationship between these two images?

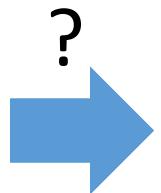


Answer: Similarity transformation (translation, rotation, uniform scale)

# Similarity?

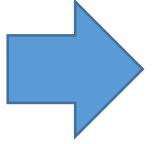
---

What is the geometric relationship between these two images?



# Similarity?

---



**Very important for creating mosaics!**

First, we need to know what this transformation is.

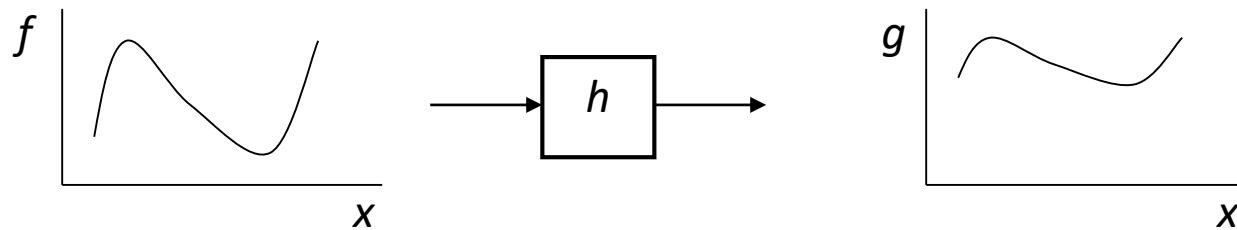
Second, we need to figure out how to compute it using feature matches.

# Image Warping

---

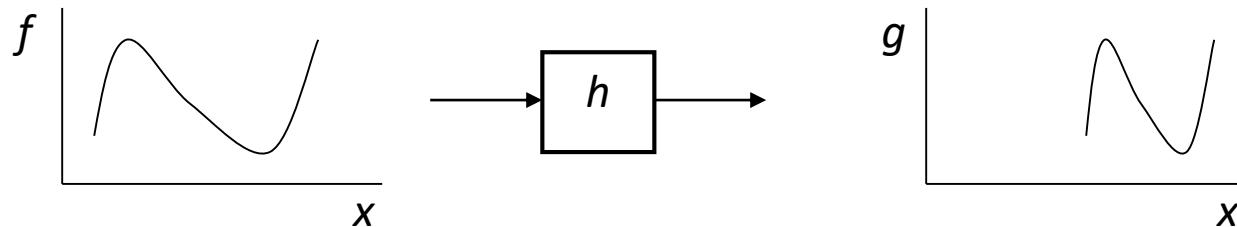
- image filtering: change *range* of image

$$g(x) = h(f(x))$$



- image warping: change *domain* of image

$$g(x) = f(h(x))$$

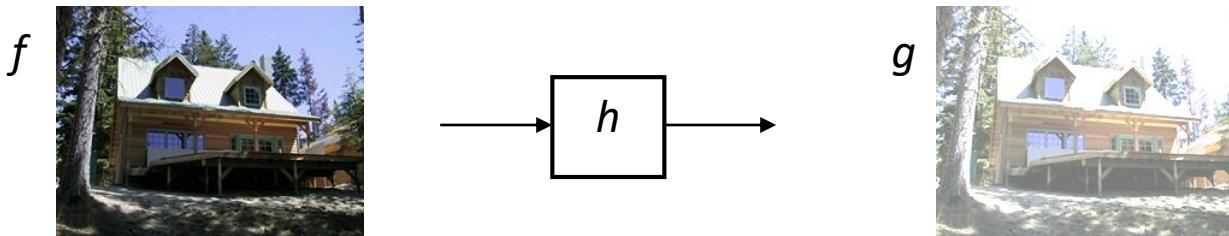


# Image Warping

---

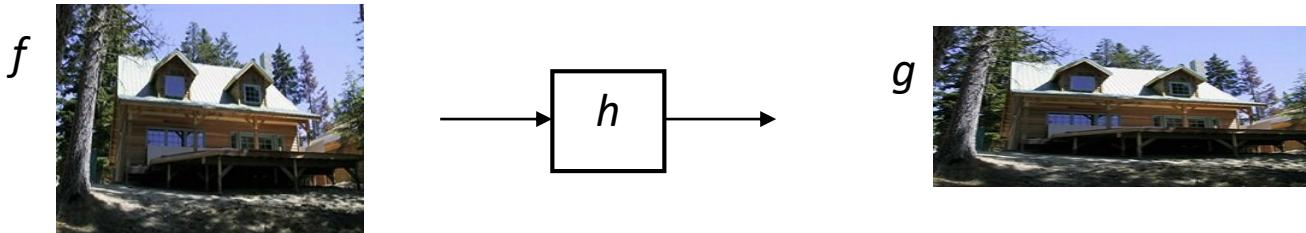
- image filtering: change *range* of image

$$g(x) = h(f(x))$$



- image warping: change *domain* of image

$$g(x) = f(h(x))$$



# Parametric (global) warping

---

Examples of parametric warps:



translation



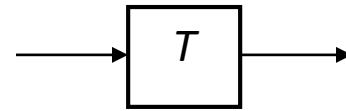
rotation



aspect

# Parametric (global) warping

---



$$\mathbf{p} = (x, y)$$

$$\mathbf{p}' = (x', y')$$

Transformation  $T$  is a coordinate-changing machine:

$$\mathbf{p}' = T(\mathbf{p})$$

What does it mean that  $T$  is **global**?

- is the same for any point  $\mathbf{p}$
- can be described by just a few numbers (parameters)

# Linear transforms

---

Let's consider *linear* transforms  
(can be represented by a 2x2 matrix):

$$\mathbf{p}' = \mathbf{T}\mathbf{p}$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \mathbf{T} \begin{bmatrix} x \\ y \end{bmatrix}$$

# Scaling

---

Uniform scaling by  $s$ :

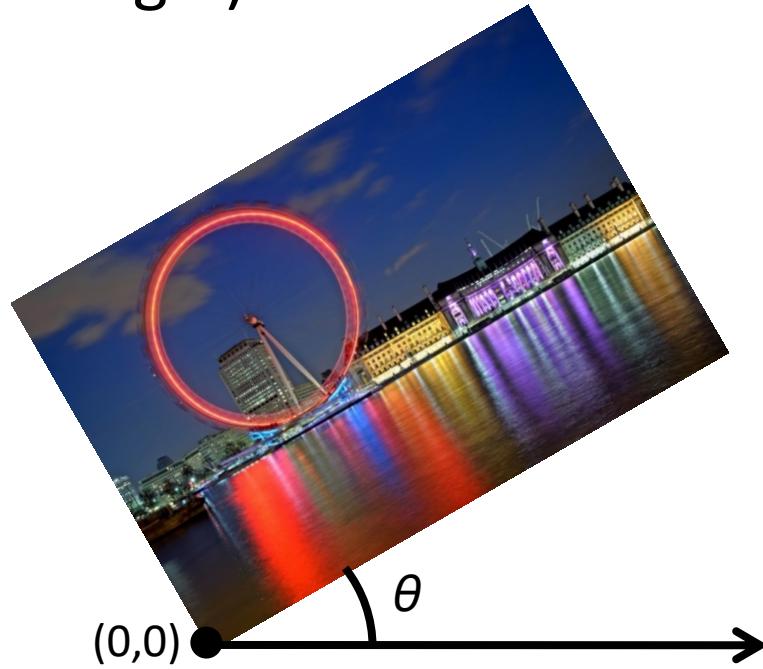


$$S = \begin{bmatrix} s & 0 \\ 0 & s \end{bmatrix}$$

# Rotation

---

Rotation by angle  $\theta$  (about the origin)



$$\mathbf{R} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

# 2x2 Matrices

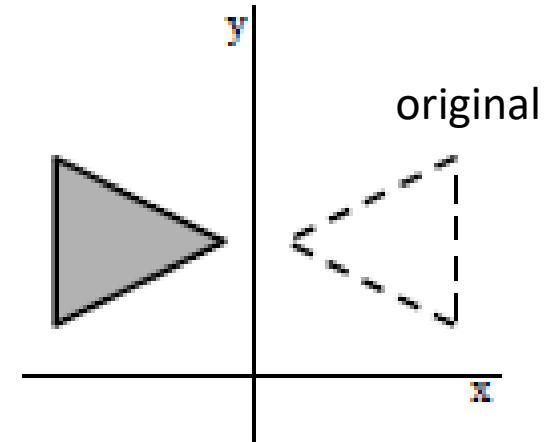
---

What types of transformations can be represented with a 2x2 matrix?

2D mirror about Y axis?

$$\begin{aligned}x' &= -x \\y' &= y\end{aligned}$$

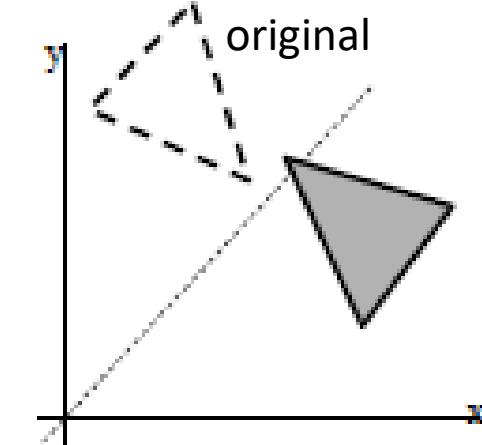
$$T = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$



2D mirror across line  $y = x$ ?

$$\begin{aligned}x' &= y \\y' &= x\end{aligned}$$

$$T = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$



# 2x2 Matrices

---

What types of transformations can be represented with a 2x2 matrix?

2D Translation?

$$x' = x + t_x$$

$$y' = y + t_y$$

NO!

Translation is not a linear operation on 2D coordinates

# All 2D Linear Transformations

---

Linear transformations are combinations of ...

- Scale,
- Rotation,
- Shear, and
- Mirror

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

Properties of linear transformations:

- Origin maps to origin
- Lines map to lines
- Parallel lines remain parallel
- Ratios are preserved
- Closed under composition

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} e & f \\ g & h \end{bmatrix} \begin{bmatrix} i & j \\ k & l \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

# Homogeneous Coordinates

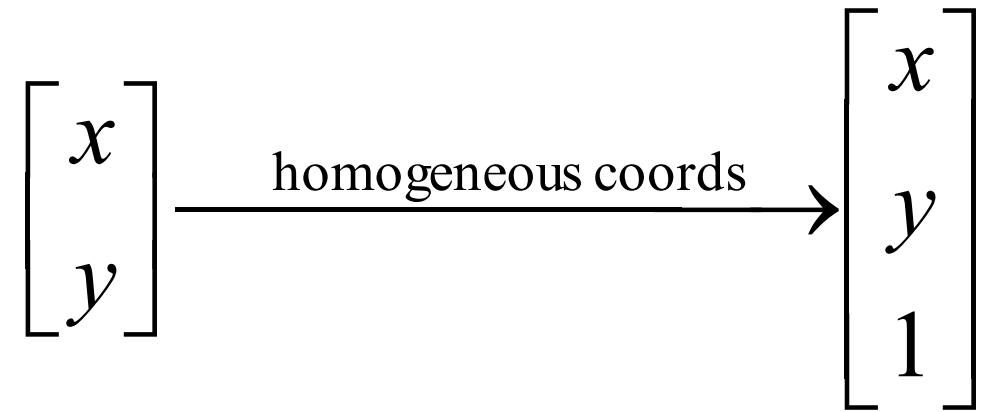
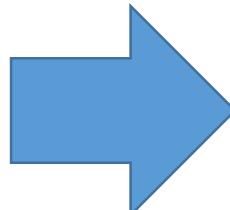
---

How can we represent  
translation

$$x' = x + t_x$$

$$y' = y + t_y$$

as a 3x3 matrix?



represent coordinates  
in 2 dimensions with a  
3-vector

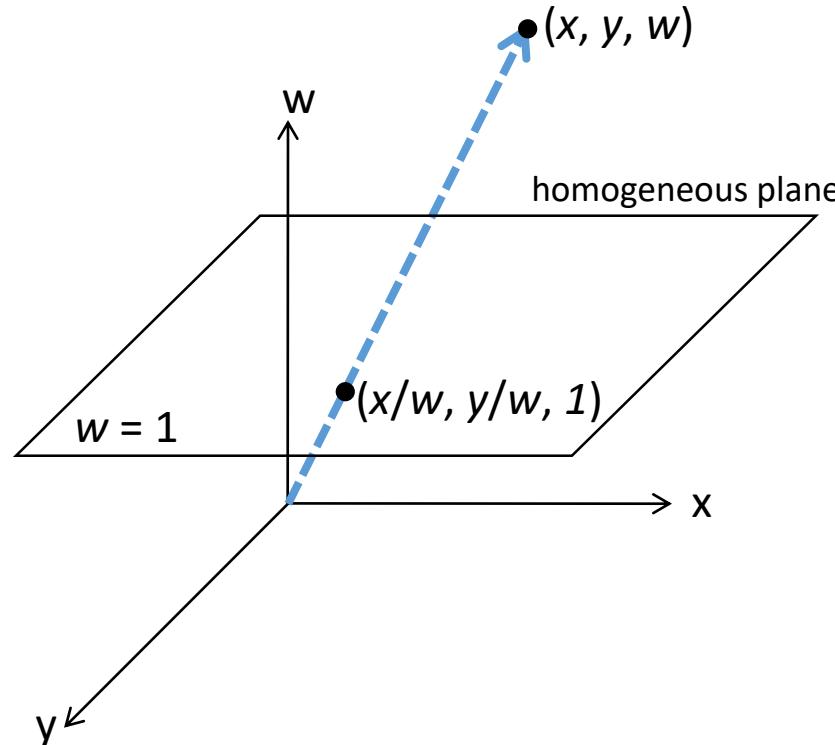
# Homogeneous coordinates

---

Trick: add one more coordinate:

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image  
coordinates



Converting *from* homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

# Translation

---

Solution using homogeneous coordinates

$$\mathbf{T} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \\ 1 \end{bmatrix}$$

# Affine transformations

---

$$\mathbf{T} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix}$$

any transformation represented by  
a 3x3 matrix with last row [ 0 0 1 ]  
we call an *affine transformation*

$$\begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix}$$

# Affine transformations

---

Affine transformations are combinations of ...

- Linear transformations, and
- Translations

$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Properties of affine transformations:

- Origin does not necessarily map to origin
- Lines map to lines
- Parallel lines remain parallel
- Ratios are preserved
- Closed under composition

# Basic affine transformations

---

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

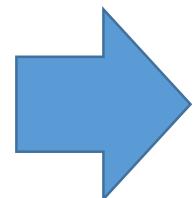
2D *in-plane* rotation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear

# É una trasformazione affine?

---

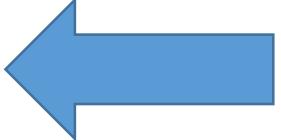


# Trasformazioni non affini

---

$$\begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix}$$

affine  
transformation

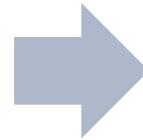


Cosa accade  
se cambiamo  
gli elementi  
della terza  
riga?

# Omografie

---

$$\mathbf{H} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix}$$

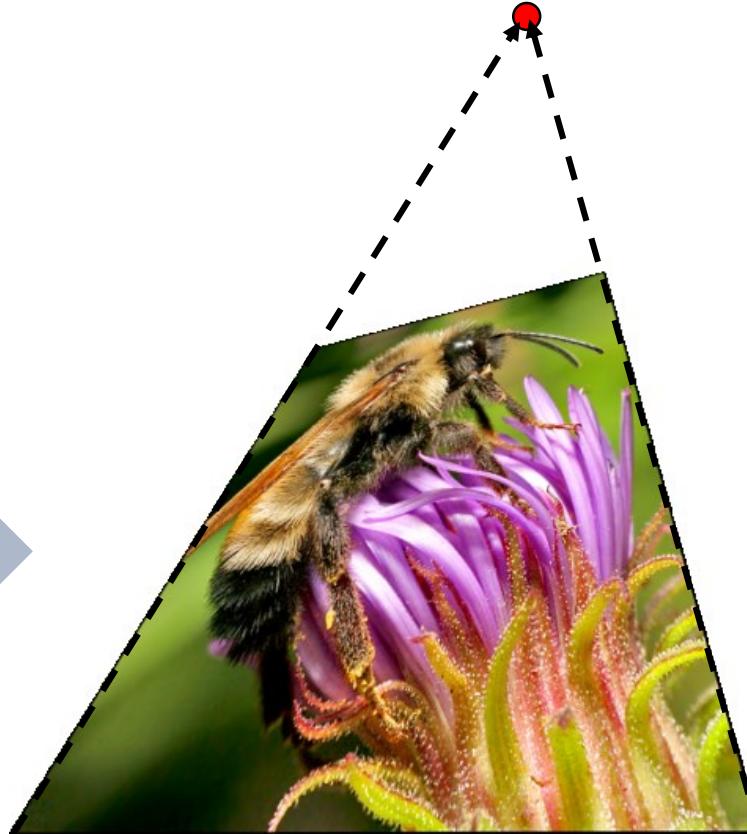
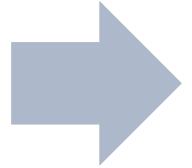
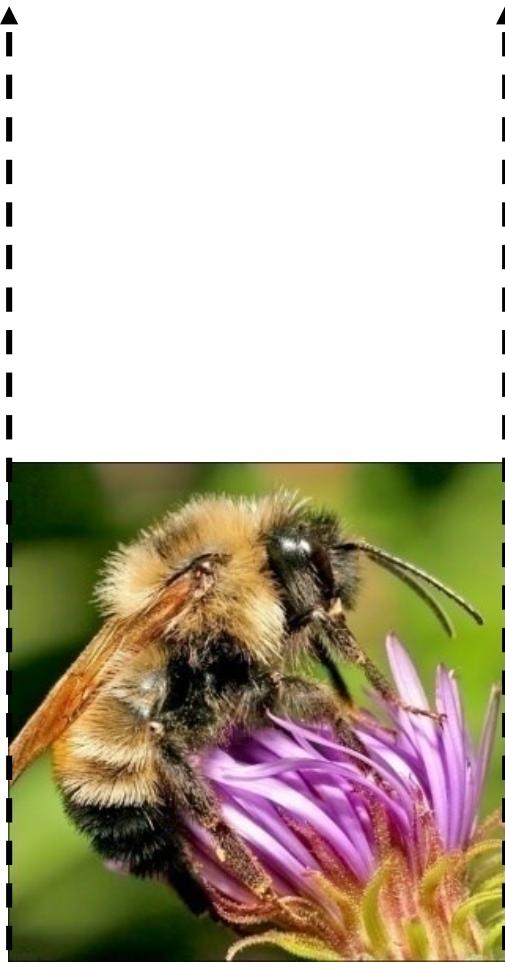


$\mathbf{H}$  is a *homography*  
(or *planar perspective map*)

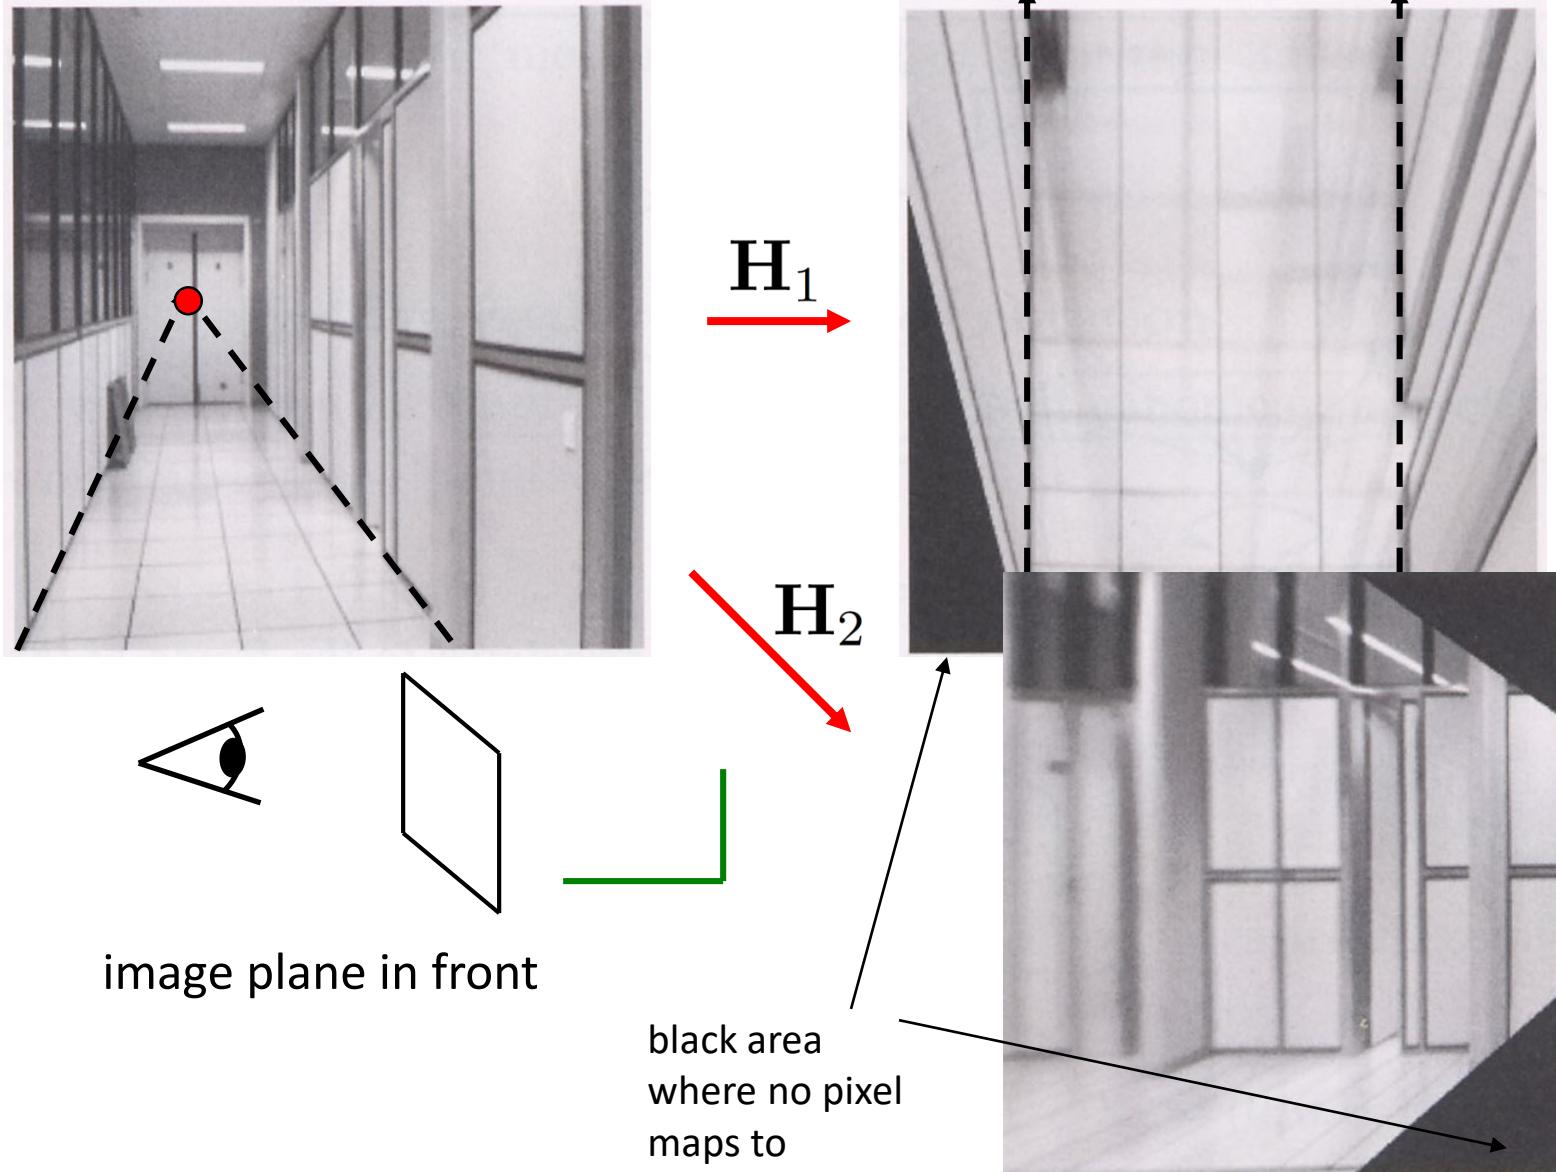


# Punti all'infinito

---

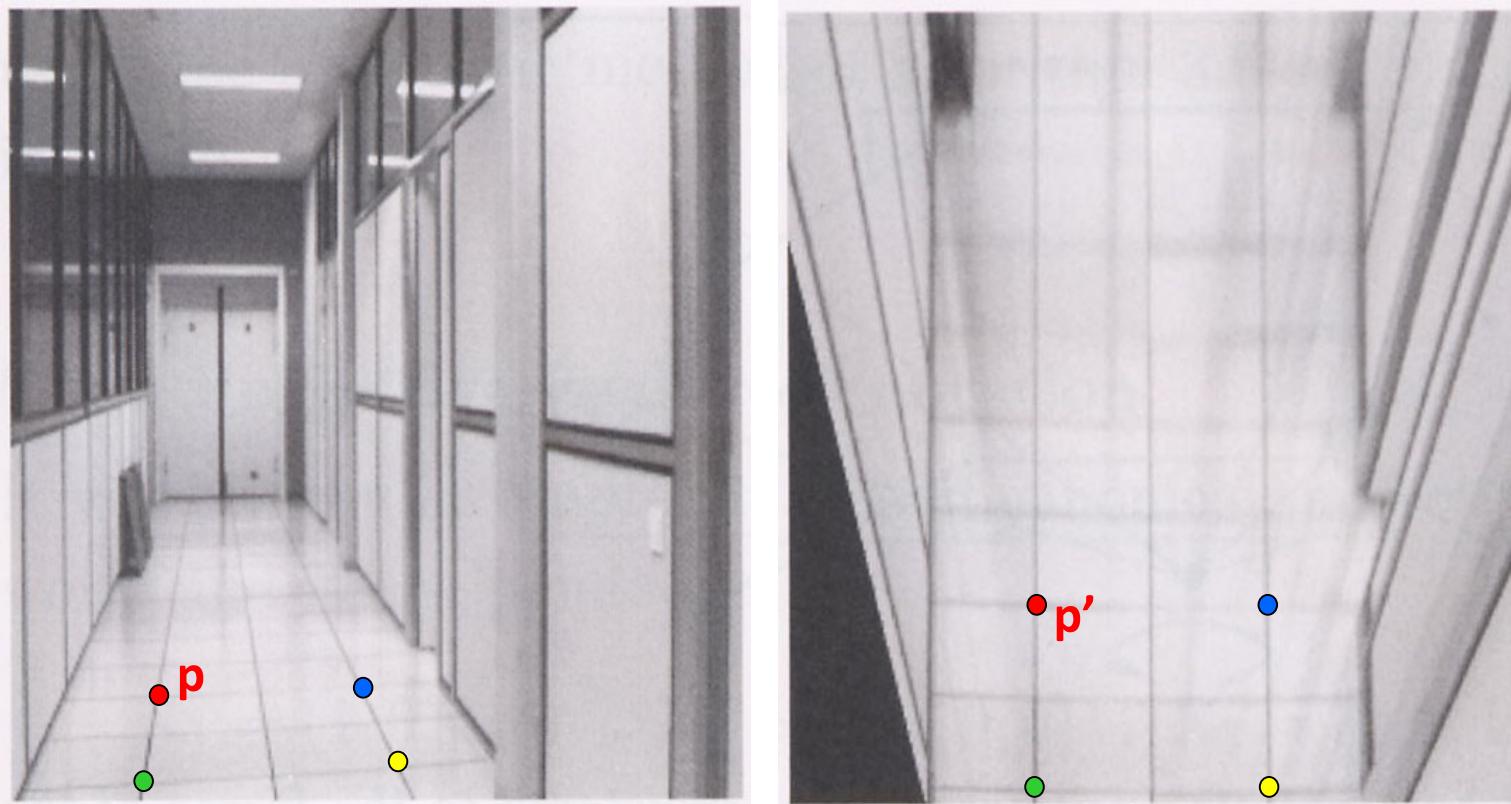


# Image warping con omografie

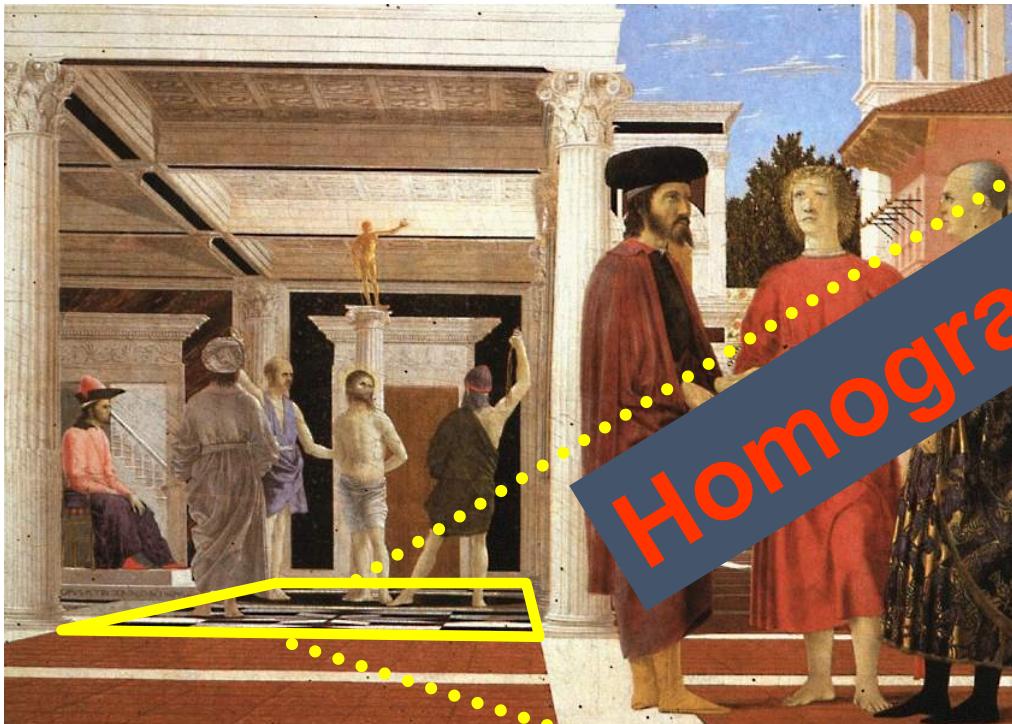


# Image warping con omografie

---



# What is the shape of the b/w floor pattern?



Homography



The floor (enlarged)

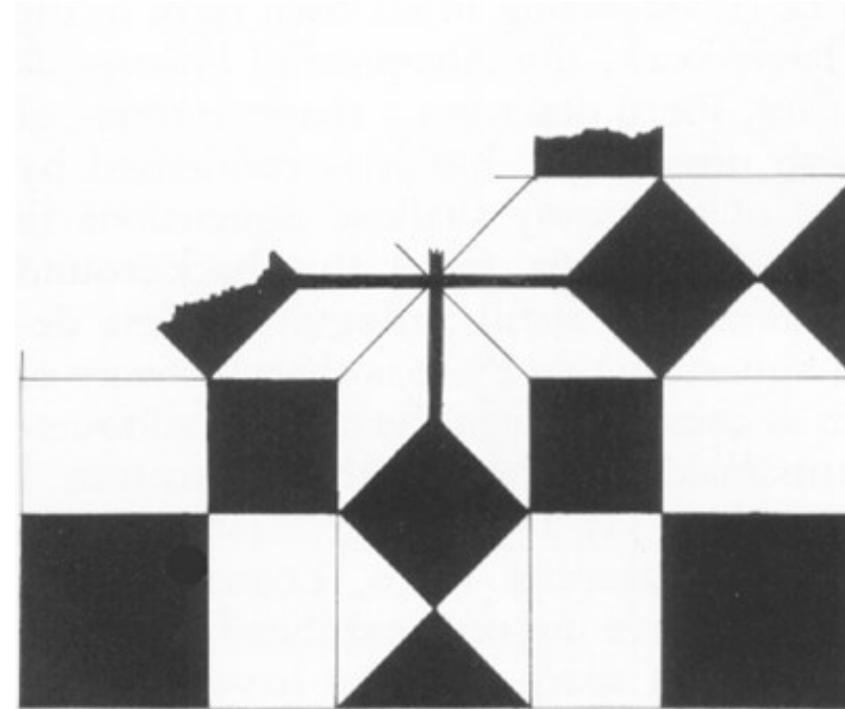


Automatically  
rectified floor

# Analyzing patterns and shapes

---

Automatic rectification



From Martin Kemp *The Science of Art*  
*(manual reconstruction)*

2 patterns have been discovered !

# Analyzing patterns and shapes

---



***St. Lucy Altarpiece, D. Veneziano***

What is the (complicated) shape of the floor pattern?



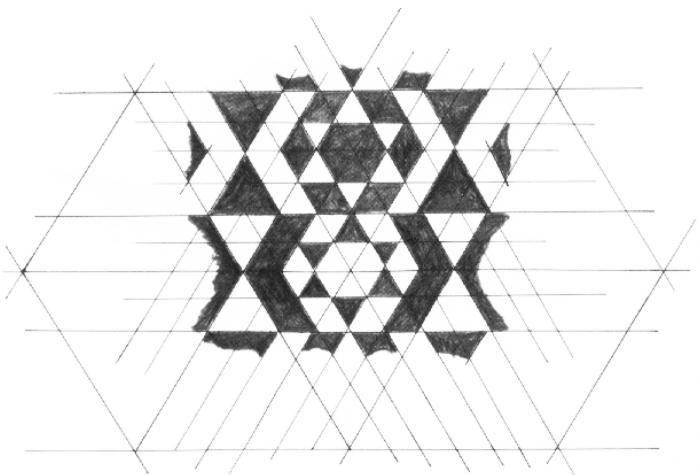
Automatically rectified floor

# Analyzing patterns and shapes

---



**Automatic  
rectification**



**From Martin Kemp, *The Science of Art*  
(manual reconstruction)**

# Analyzing patterns and shapes

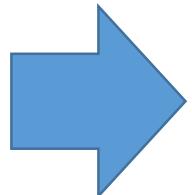
---



The Ambassadors by Hans Holbein the Younger, 1533

# É una omografia

---



# Omografie

---

Homographies ...

- Affine transformations, and
- Projective warps

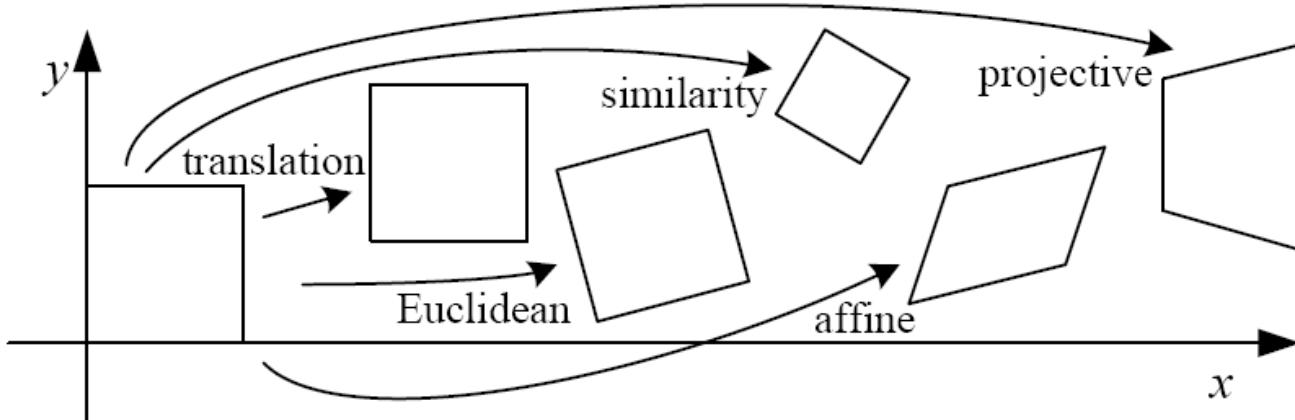
$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Properties of projective transformations:

- Origin does not necessarily map to origin
- Lines map to lines
- Parallel lines do not necessarily remain parallel
- Ratios are not preserved
- Closed under composition

# Ricapitolando

---

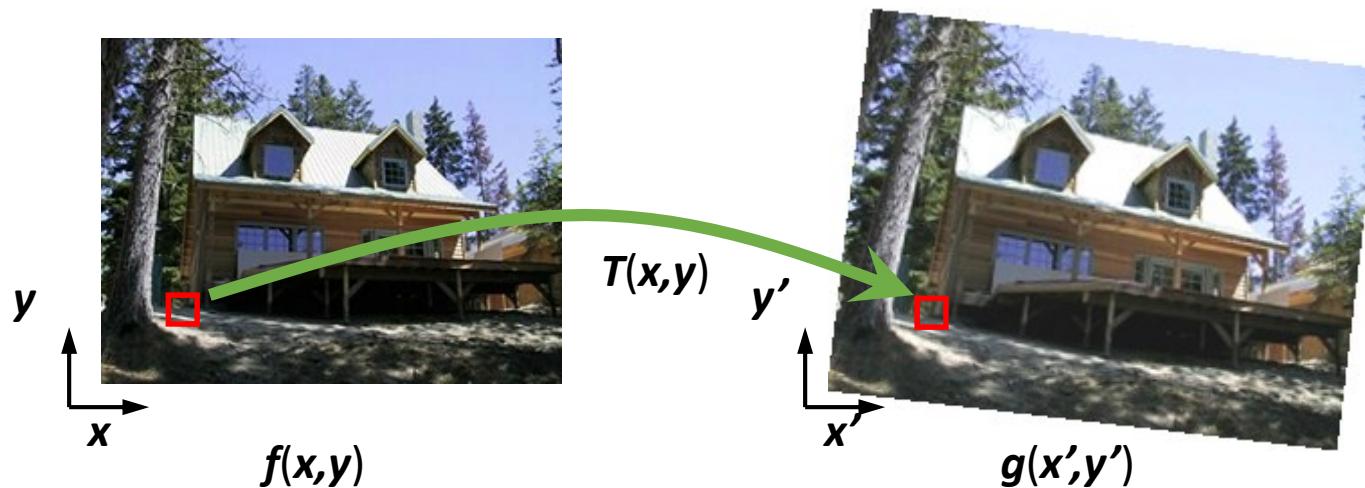


Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

# Implementing image warping

---

Given a coordinate transform  $(x',y') = T(x,y)$  and a source image  $f(x,y)$ , how do we compute a transformed image  $g(x',y') = f(T(x,y))$ ?

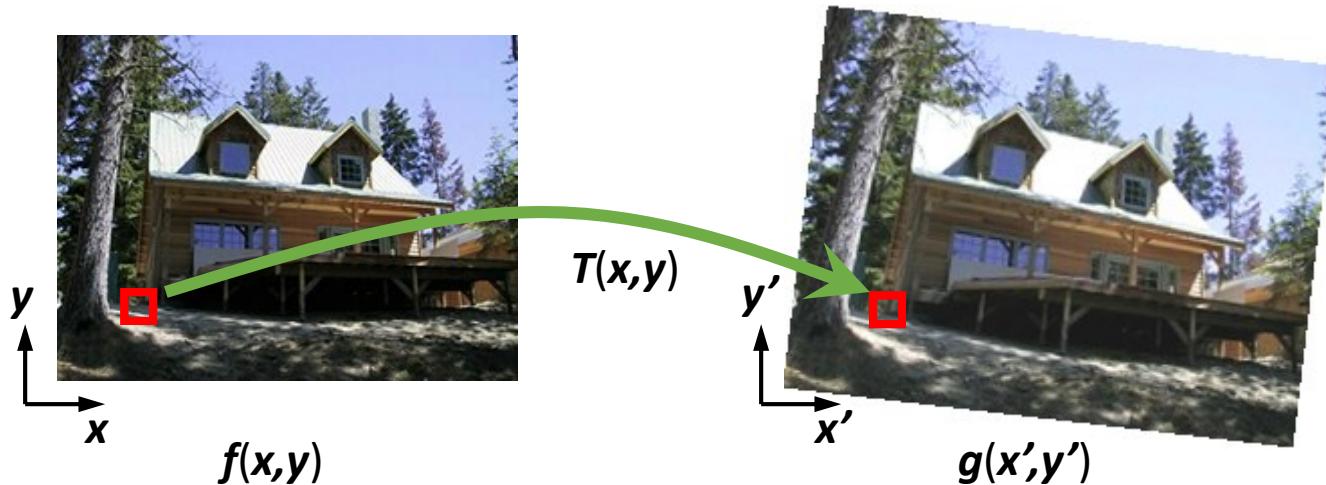


# Forward Warping

---

Send each pixel  $f(x)$  to its corresponding location  
 $(x',y') = T(x,y)$  in  $g(x',y')$

- What if pixel lands “between” two pixels?

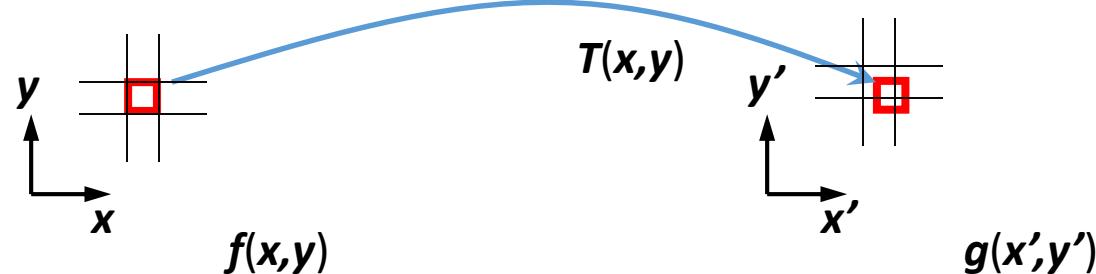


# Forward Warping

---

Send each pixel  $f(x)$  to its corresponding location  
 $(x',y') = T(x,y)$  in  $g(x',y')$

- What if pixel lands “between” two pixels?
- Answer: add “contribution” to several pixels, normalize later (splatting)
- Can still result in holes

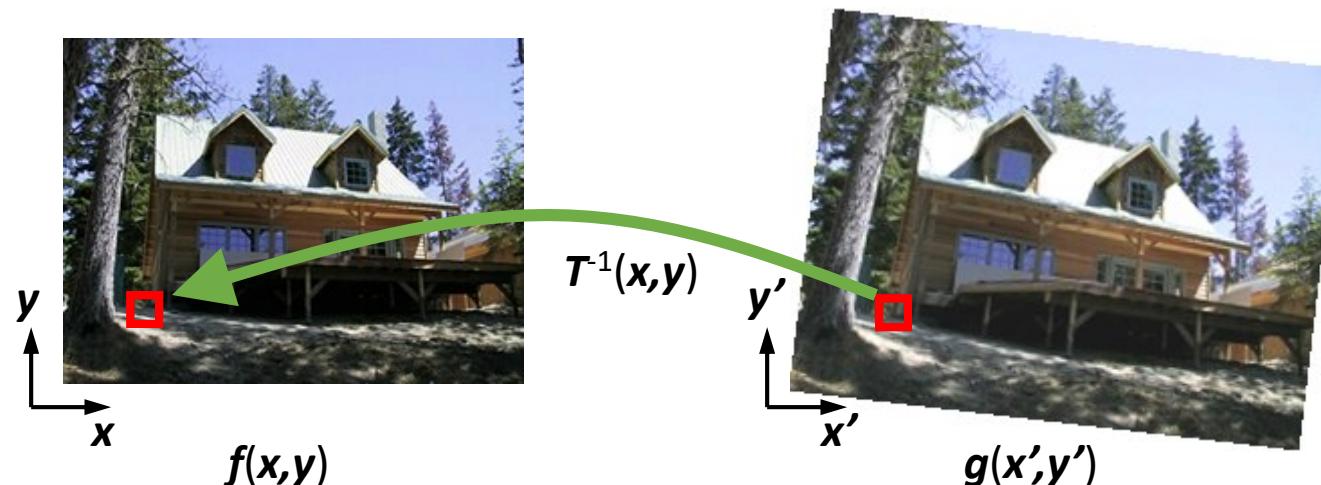


# Inverse Warping

---

Get each pixel  $g(x',y')$  from its corresponding location  
 $(x,y) = T^{-1}(x',y')$  in  $f(x,y)$

- Requires taking the inverse of the transform
- What if pixel comes from “between” two pixels?

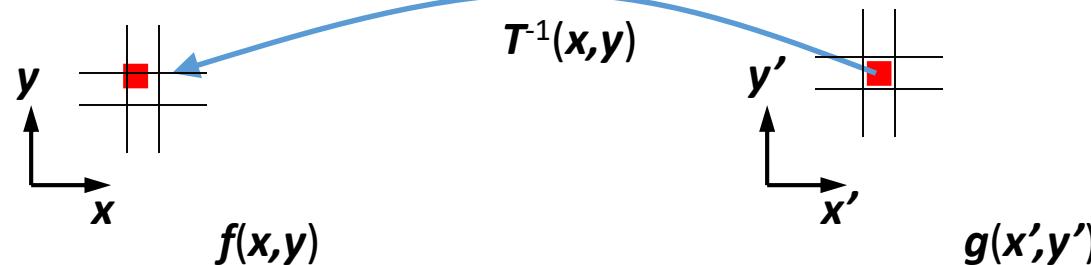


# Inverse Warping

---

Get each pixel  $g(x')$  from its corresponding location  
 $x' = h(x)$  in  $f(x)$

- What if pixel comes from “between” two pixels?
- Answer: *resample color value from interpolated (prefiltered) source image*



# Interpolation

---

Possible interpolation filters:

- nearest neighbor
- bilinear
- bicubic
- sinc

Needed to prevent “jaggies”  
and “texture crawl”  
(with prefiltering)



# Forward vs. inverse warping

---

Q: Which is better?

A: usually inverse—eliminates holes

- however, it requires an invertible warp function—not always possible...

# Esempio omografia+warping



```
import cv2 as cv
from google.colab.patches import cv2_imshow
from urllib.request import urlopen
import numpy as np

req_left = urlopen('https://dbloisi.github.io/corsi/images/montagna-1.jpg')
arr_left = np.array(bytarray(req_left.read()), dtype=np.uint8)
img_left = cv.imdecode(arr_left, -1)
cv2_imshow(img_left)
```



# Esempio omografia+warping

```
▶ req_right = urlopen('https://dbloisi.github.io/corsi/images/montagna-2.jpg')
arr_right = np.array(bytearray(req_right.read()), dtype=np.uint8)
img_right = cv.imdecode(arr_right, -1)
cv2.imshow(img_right)
```



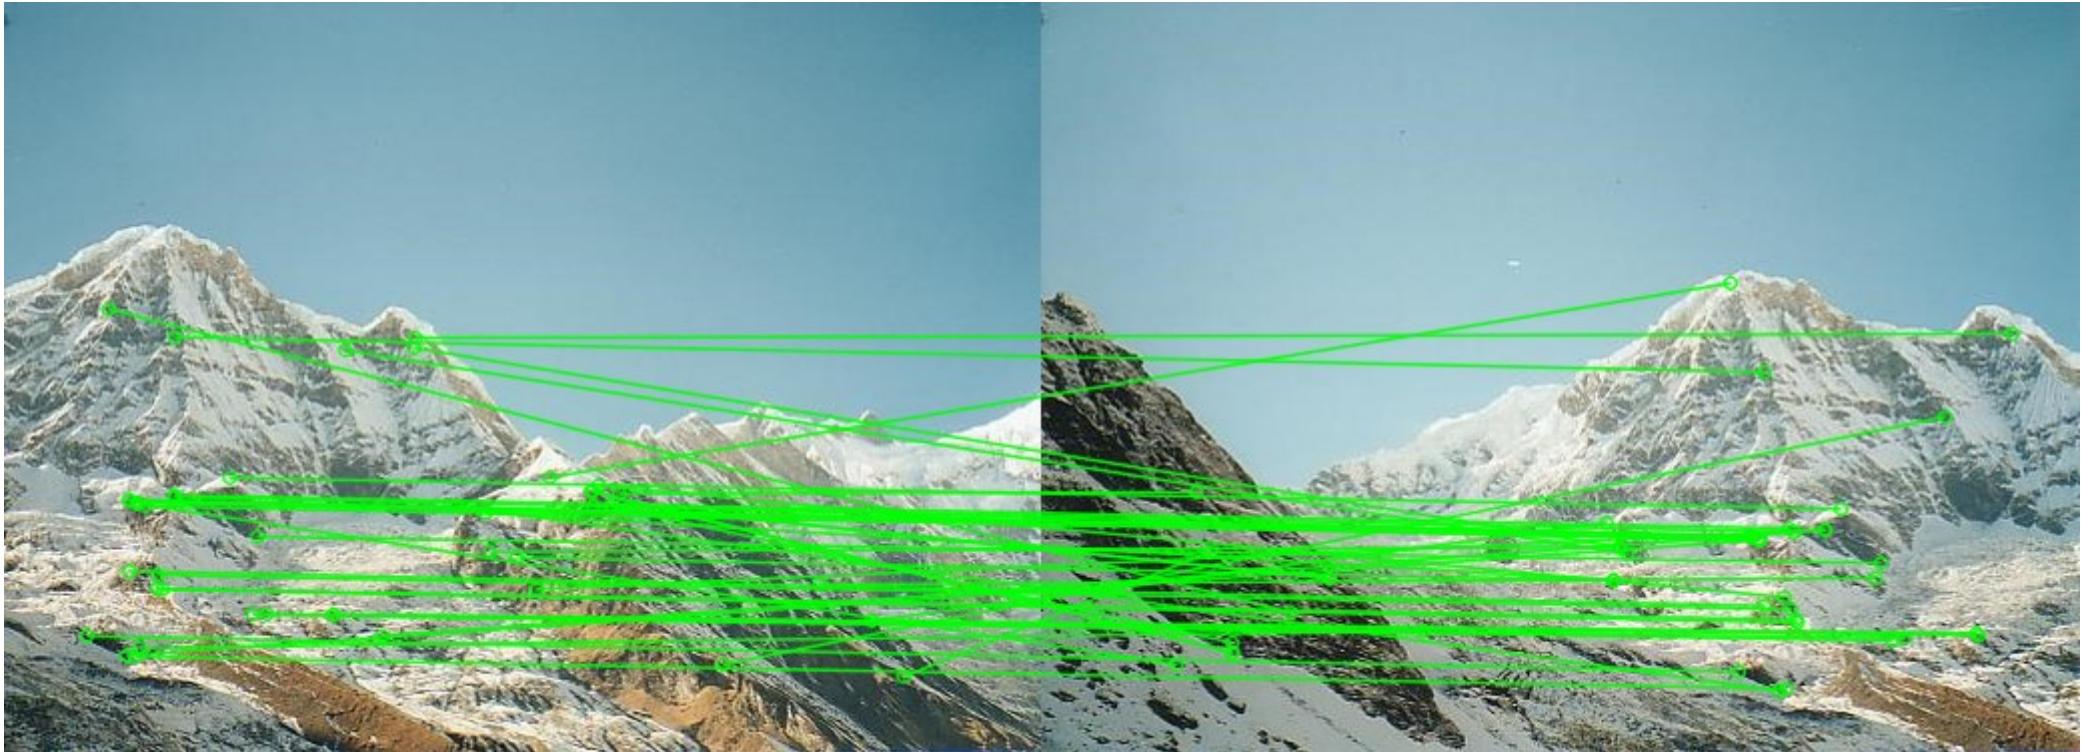
# Esempio omografia+warping

```
# orb descriptor
orb = cv.ORB_create()
# find key points
kp1, des1 = orb.detectAndCompute(img_right, None)
kp2, des2 = orb.detectAndCompute(img_left, None)
# brute force matching
match = cv.BFMatcher()
matches = match.knnMatch(des1,des2,k=2)
# distance ratio
# "Distinctive Image Features from Scale-Invariant Keypoints"
# by David G. Lowe
good = []
for m,n in matches:
    if m.distance < 0.85*n.distance:
        good.append(m)

# drawing good matches
draw_params = dict(matchColor=(0,255,0),
                    singlePointColor=None,
                    flags=2)
matches_img = cv.drawMatches(img_right,kp1,img_left,kp2,good,None,**draw_params)
cv2_imshow(matches_img)
```

# Esempio omografia+warping

---



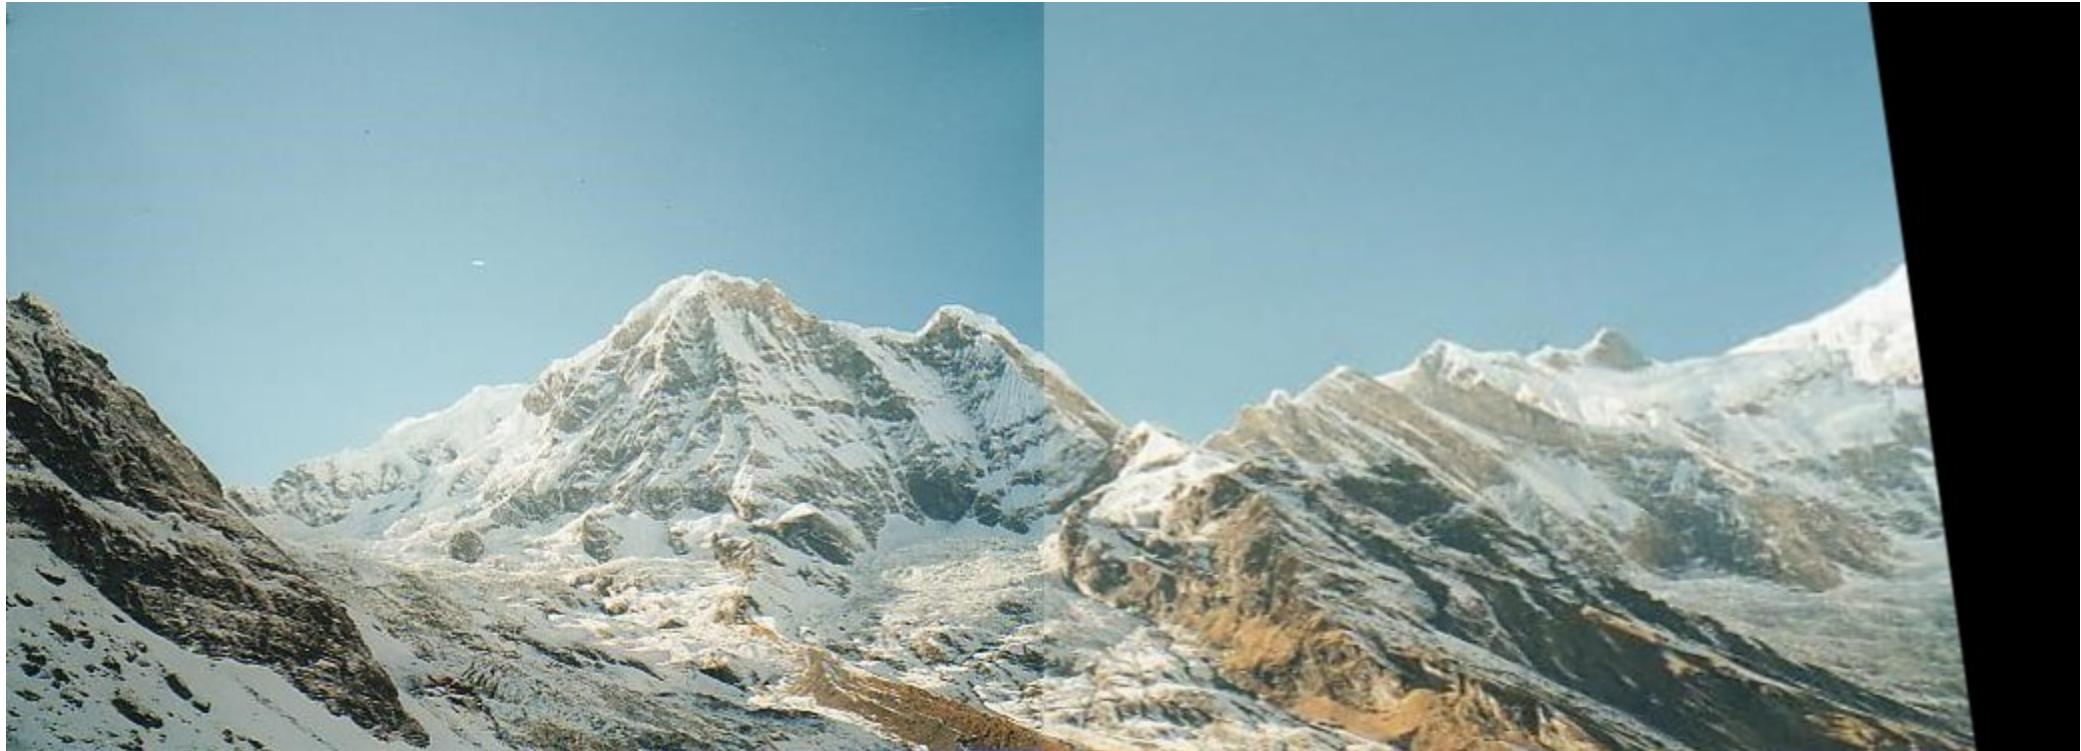
# Esempio omografia+warping

```
# homography computation
MIN_MATCH_COUNT = 5
if len(good) > MIN_MATCH_COUNT:
    src_pts = np.float32([ kp1[m.queryIdx].pt for m in good ]).reshape(-1,1,2)
    dst_pts = np.float32([ kp2[m.trainIdx].pt for m in good ]).reshape(-1,1,2)
    M, mask = cv.findHomography(src_pts, dst_pts, cv.RANSAC, 5.0)
    h,w,c = img_right.shape
    pts = np.float32([ [0,0],[0,h-1],[w-1,h-1],[w-1,0] ]).reshape(-1,1,2)
    dst = cv.perspectiveTransform(pts, M)
else:
    print("Not enought matches are found - %d/%d", (len(good)/MIN_MATCH_COUNT))
```

# Esempio omografia+warping

---

```
# creating panorama image
panorama_img = cv.warpPerspective(img_right,M,(img_left.shape[1] + img_right.shape[1], img_left.shape[0]))
panorama_img[0:img_left.shape[0],0:img_left.shape[1]] = img_left
cv2_imshow(panorama_img)
```

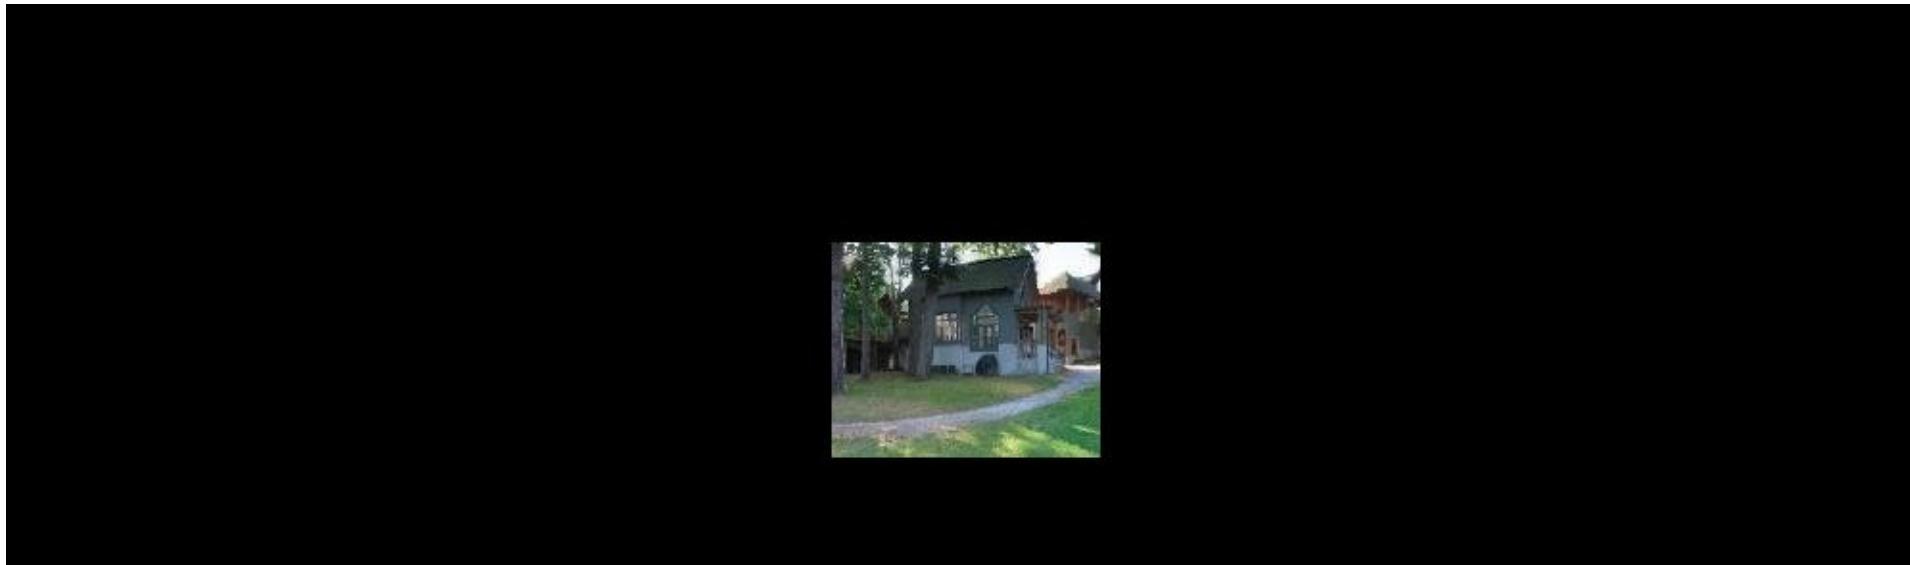


# Panoramas

---

Are you getting the whole picture?

Compact Camera FOV =  $50 \times 35^\circ$



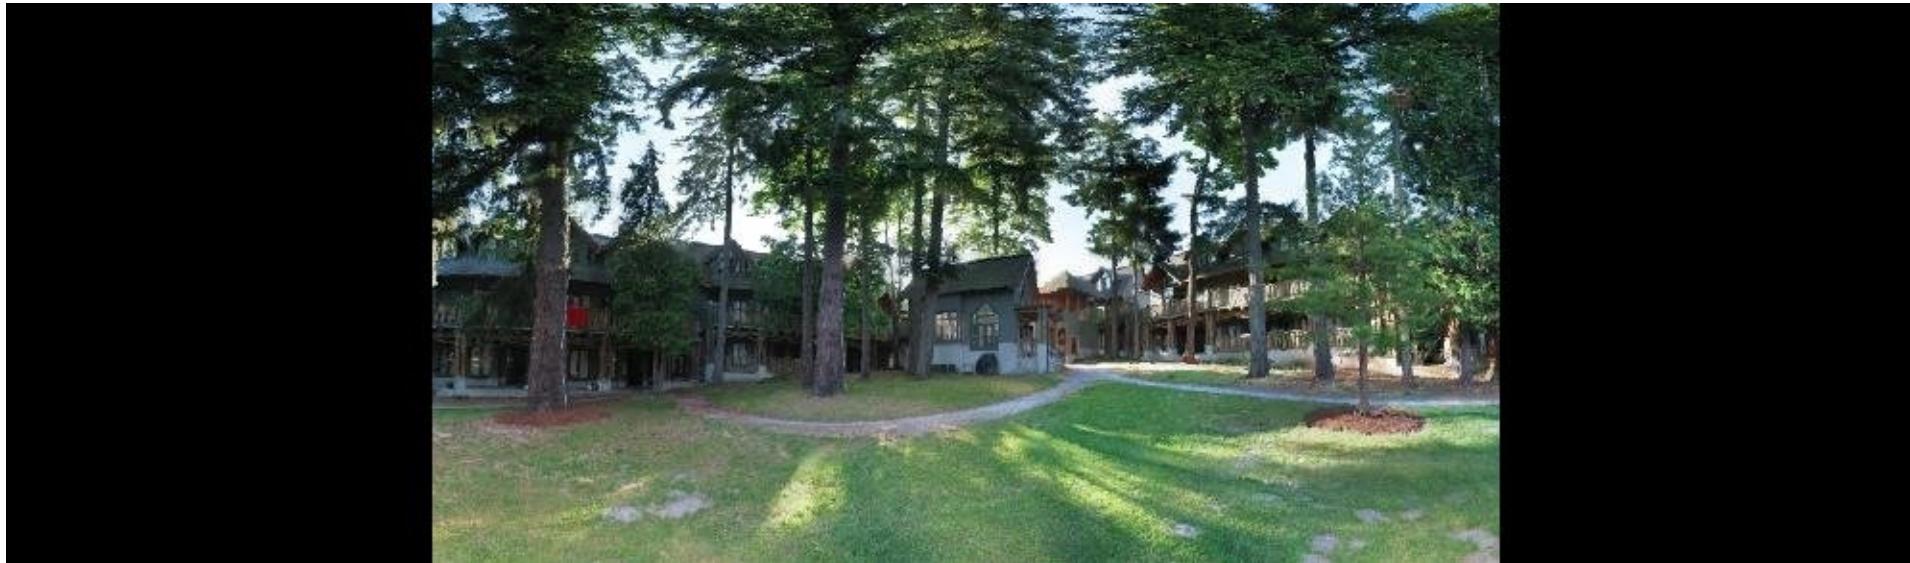
# Panoramas

---

Are you getting the whole picture?

Compact Camera FOV =  $50 \times 35^\circ$

Human FOV =  $200 \times 135^\circ$



# Panoramas

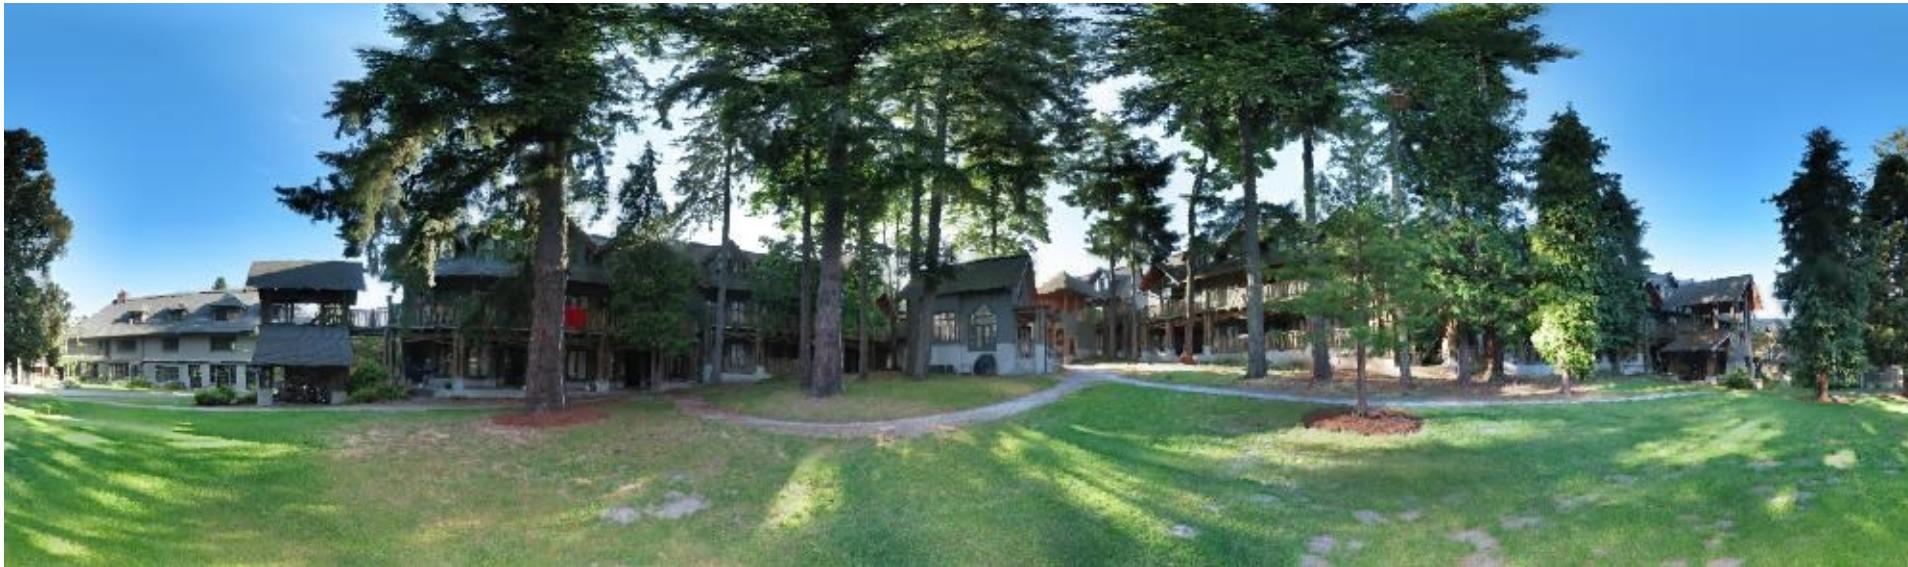
---

Are you getting the whole picture?

Compact Camera FOV =  $50 \times 35^\circ$

Human FOV =  $200 \times 135^\circ$

Panoramic Mosaic =  $360 \times 180^\circ$



# Why “Recognising Panoramas”?

---

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



# Why “Recognising Panoramas”?

---

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



# Why “Recognising Panoramas”?

---

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



- 2D Rotations ( $\theta, \phi$ )
  - Ordering  $\not\Rightarrow$  matching images

# Why “Recognising Panoramas”?

---

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images



- 2D Rotations ( $\theta, \phi$ )
  - Ordering  $\not\Rightarrow$  matching images



# Why “Recognising Panoramas”?

---

- 1D Rotations ( $\theta$ )
  - Ordering  $\Rightarrow$  matching images

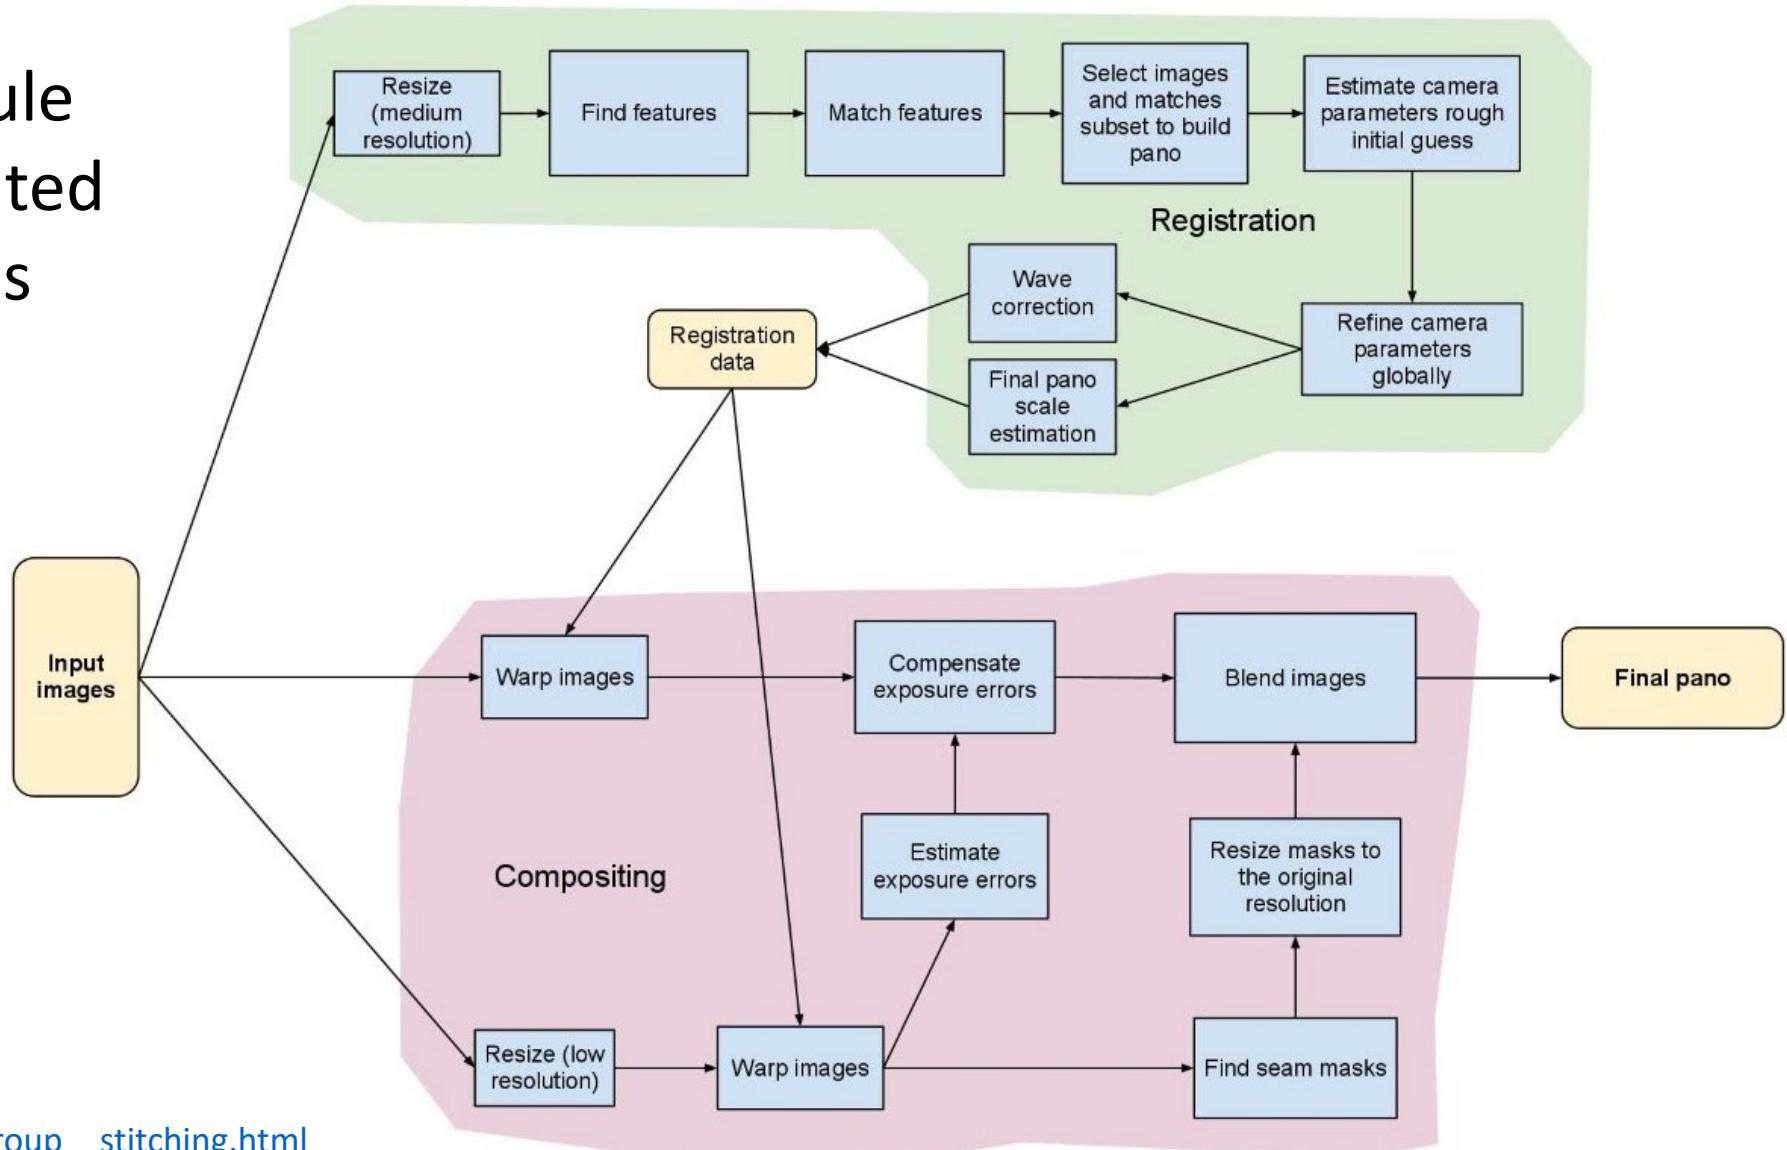


- 2D Rotations ( $\theta, \phi$ )
  - Ordering  $\not\Rightarrow$  matching images



# Stitching in OpenCV

The stitching module pipeline implemented in the [Stitcher](#) class



# Stitching in OpenCV

---



# Stitching in OpenCV

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** CO stitching.ipynb ☆
- Menu Bar:** File Edit View Insert Runtime Tools Help
- Left Sidebar (Files):** Shows a file tree with a folder named "sample\_data" containing five files: Univ1.jpg, Univ2.jpg, Univ3.jpg, Univ4.jpg, and Univ5.jpg. It also includes "Upload", "Refresh", and "Mount Drive" buttons.
- Code Cell:** Contains Python code for reading input images and appending them to a list. A play button icon is present in the cell header.

```
import numpy as np
import cv2 as cv
import sys

# read input images
imgs = []
for i in range(1,6):
    img_name = "Univ" + str(i) + ".jpg"
    img = cv.imread(img_name)
    if img is None:
        print("can't read image " + img_name)
        sys.exit(-1)
    else:
        print(img_name + " loaded")
        imgs.append(img)
```

# Stitching in OpenCV

---

```
stitcher = cv.Stitcher.create(cv.Stitcher_PANORAMA)

status, pano = stitcher.stitch(imgs)

if status != cv.Stitcher_OK:
    print("Can't stitch images, error code = %d" % status)
    sys.exit(-1)

cv.imwrite("panorama.jpg", pano)

print("stitching completed successfully.")
```

⇨ stitching completed successfully.

# Stitching in OpenCV

---



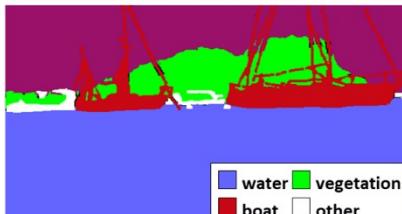
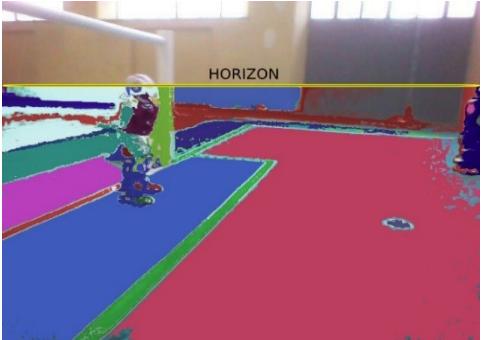


**UNIVERSITÀ DEGLI STUDI  
DELLA BASILICATA**

*Corso di Visione e Percezione*  
A.A. 2019/2020

# Omografie

Aprile 2020



Docente  
**Domenico Daniele Bloisi**